

Chapter 9 – Uniprocessor Scheduling

Lecture 8

Roadmap


- **Types of Processor Scheduling**
- **Scheduling Algorithms**
- Traditional UNIX Scheduling

Scheduling

- An OS must allocate resources amongst competing processes.
- The resource provided by a processor is execution time




Sched



Max Ticks Allowed = 15

Scheduling Objectives



- The scheduling function should
 - Share time *fairly* among processes
 - Prevent starvation of a process
 - Use the processor efficiently
 - Have low overhead
 - Prioritise processes when necessary (e.g. real time deadlines)


Types of Scheduling

Table 9.1 Types of Scheduling

| | |
|-------------------------------|--|
| Long-term scheduling | The decision to add to the pool of processes to be executed |
| Medium-term scheduling | The decision to add to the number of processes that are partially or fully in main memory |
| Short-term scheduling | The decision as to which available process will be executed by the processor |
| I/O scheduling | The decision as to which process's pending I/O request shall be handled by an available I/O device |

5 State model

Remember this diagram
from Chapter 3



(b) With Two Suspend States

Scheduling and Process State Transitions

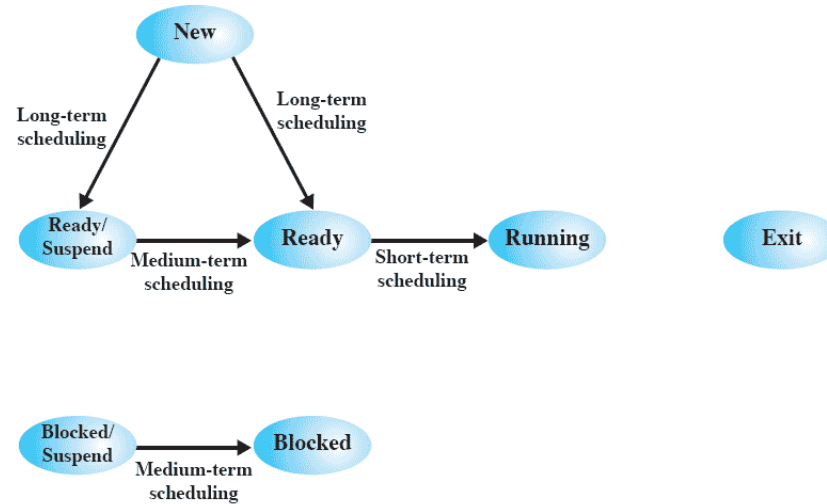
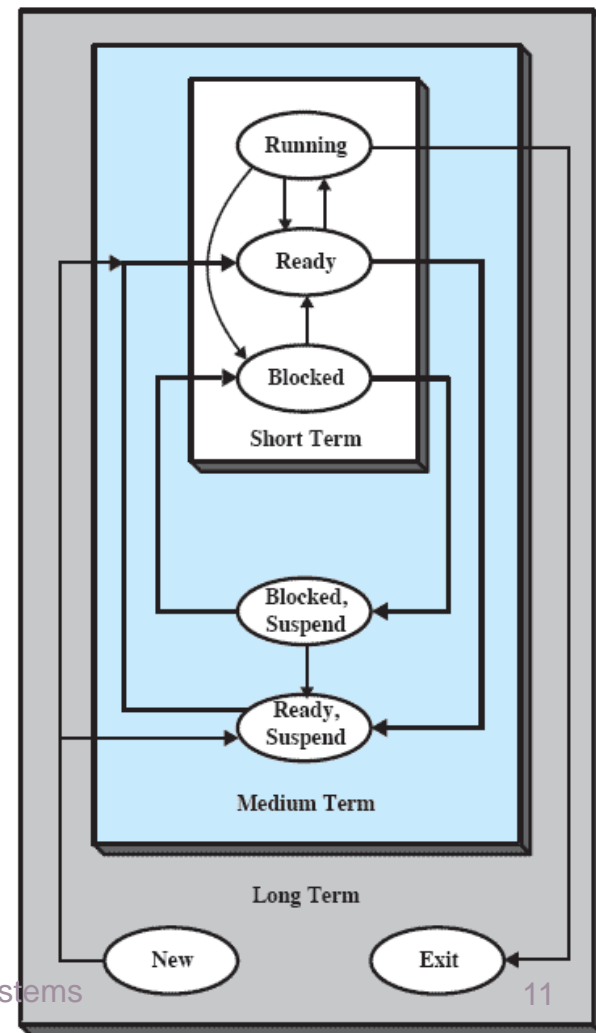
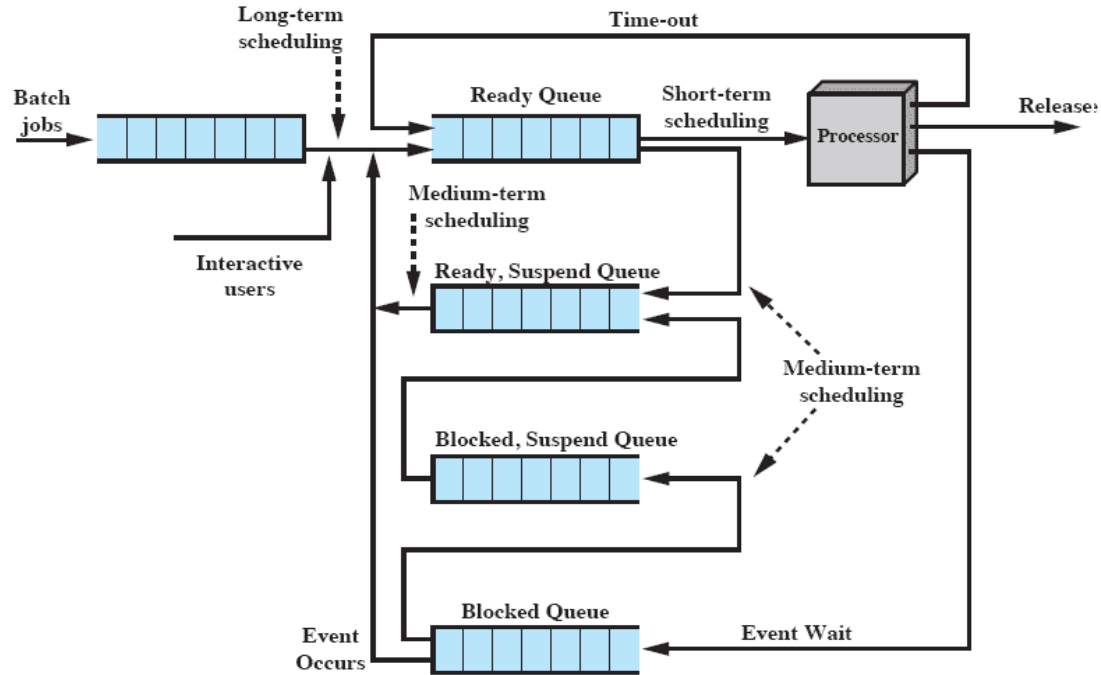


Figure 9.1 Scheduling and Process State Transitions

Nesting of Scheduling Functions



Queuing Diagram




*scheduling is
a matter of
managing
queues*


Long term scheduling




Medium Term Scheduling



Short Term Scheduling



I/O



System Calls (OS)



Signals



Timeout

Roadmap


- Types of Processor Scheduling
- **Scheduling Algorithms**
- Traditional UNIX Scheduling

Aim of Short Term Scheduling

- Main objective is to allocate processor time to optimize certain aspects of system behaviour.
 - A set of criteria is needed to evaluate the scheduling policy.

Short-Term Scheduling Criteria: User vs System


- User-oriented
- System-oriented



Short-Term Scheduling Criteria: Performance

We could differentiate between performance related criteria, and those unrelated to performance

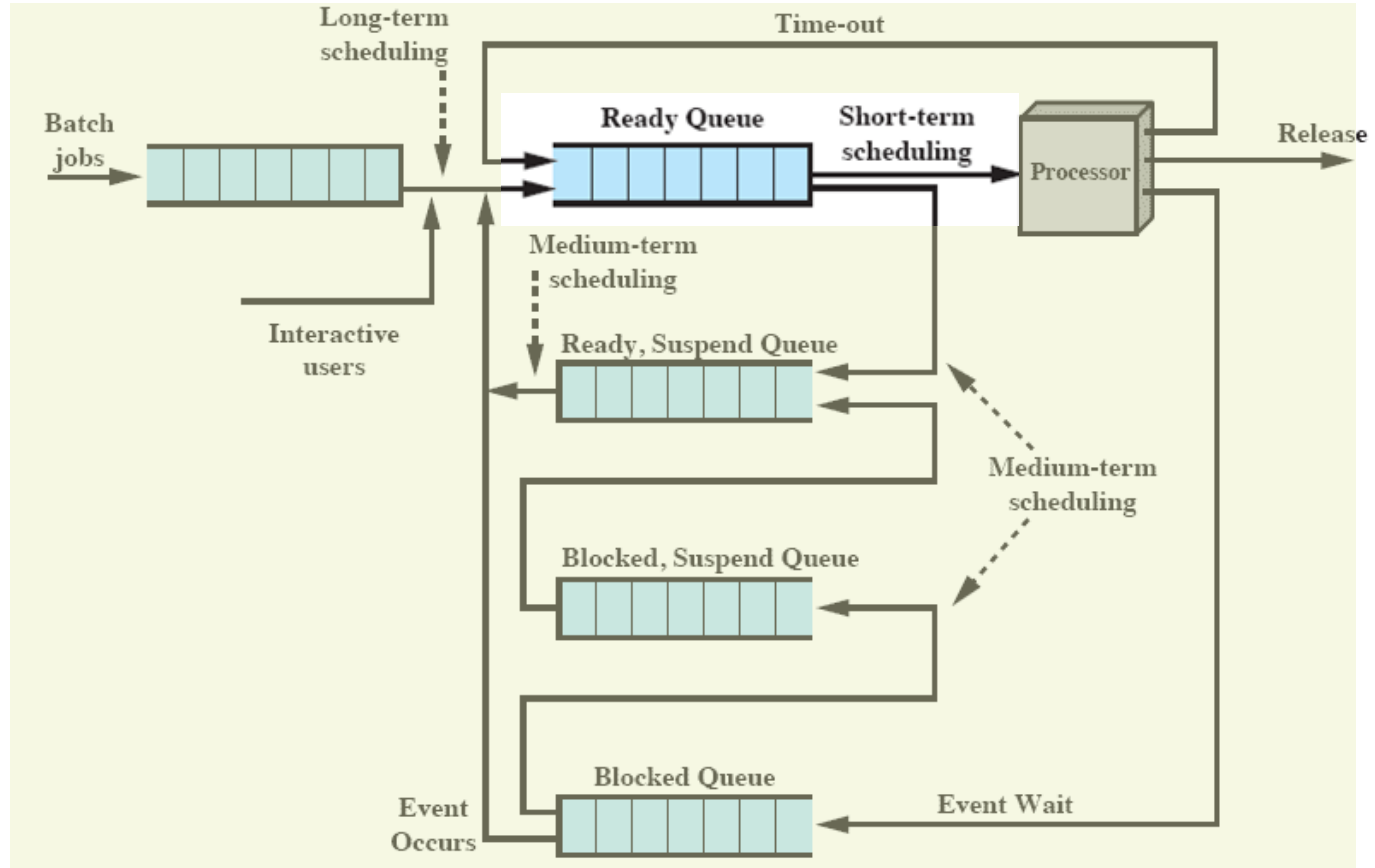
- Performance-related
- Non-performance related



Priorities

- Scheduler will always choose a process of higher priority over one of lower priority
 - Have multiple ready queues to represent each level of priority

Queuing Diagram – One ready queue



Priority Queuing

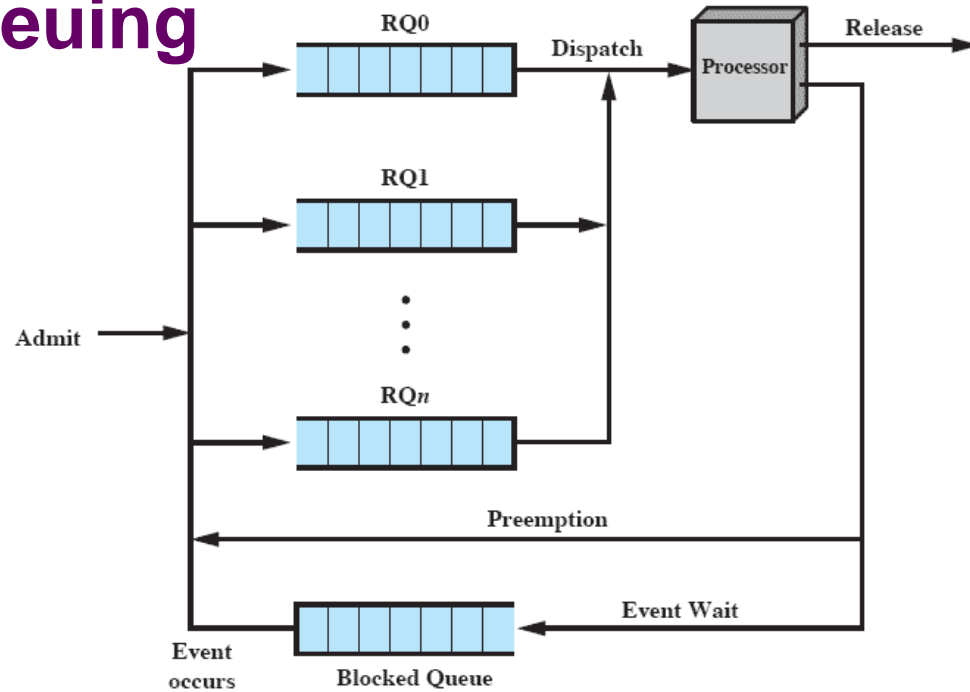


Figure 9.4 Priority Queuing

Starvation

- Problem:
 - Lower-priority may suffer starvation if there is a steady supply of high priority processes.
- Solution
 - Allow a process to change its priority based on its age or execution history

Alternative Scheduling Policies

Table 9.3 Characteristics of Various Scheduling Policies

| | FCFS | Round robin | SPN | SRT | HRRN | Feedback |
|----------------------------|---|---|---|-----------------------------|------------------------------------|-------------------------------|
| Selection function | $\max[w]$ | constant | $\min[s]$ | $\min[s - e]$ | $\max\left(\frac{w + s}{s}\right)$ | (see text) |
| Decision mode | Non-preemptive | Preemptive (at time quantum) | Non-preemptive | Preemptive (at arrival) | Non-preemptive | Preemptive (at time quantum) |
| Throughput | Not emphasized | May be low if quantum is too small | High | High | High | Not emphasized |
| Response time | May be high, especially if there is a large variance in process execution times | Provides good response time for short processes | Provides good response time for short processes | Provides good response time | Provides good response time | Not emphasized |
| Overhead | Minimum | Minimum | Can be high | Can be high | Can be high | Can be high |
| Effect on processes | Penalizes short processes; penalizes I/O bound processes | Fair treatment | Penalizes long processes | Penalizes long processes | Good balance | May favor I/O bound processes |
| Starvation | No | No | Possible | Possible | No | Possible |

Selection Function

- Determines which process is selected for execution
- If based on execution characteristics, then important quantities are:
 - w = time spent in system so far, waiting
 - e = time spent in execution so far
 - s = total service time required by the process, including e ;

Decision Mode

Non-preemptive vs Preemptive

- Non-preemptive
 - Once a process is in the running state, it will continue until it terminates or blocks itself for I/O
- Preemptive
 - Currently running process may be interrupted and moved to ready state by the OS
 - Preemption may occur when new process arrives, on an interrupt, or periodically.

Process Scheduling Example

Example set of processes, consider each a batch job

Table 9.4 Process Scheduling Example


| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| A | 0 | 3 |
| B | 2 | 6 |
| C | 4 | 4 |
| D | 6 | 5 |
| E | 8 | 2 |

Service time represents total execution time

First-Come - First-Served

- Each process joins the Ready queue
- When the current process ceases to execute, the longest (oldest) process in the Ready queue is selected


First-Come-First
Served (FCFS)



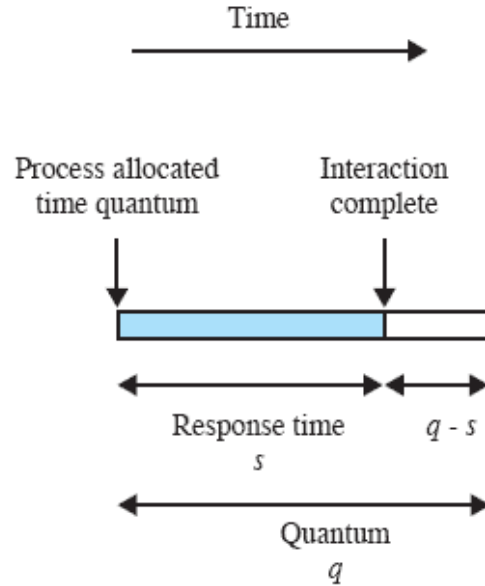
Round Robin

- Uses preemption based on a clock
 - also known as time slicing, because each process is given a slice of time before being preempted.

Round-Robin
(RR), $q = 1$

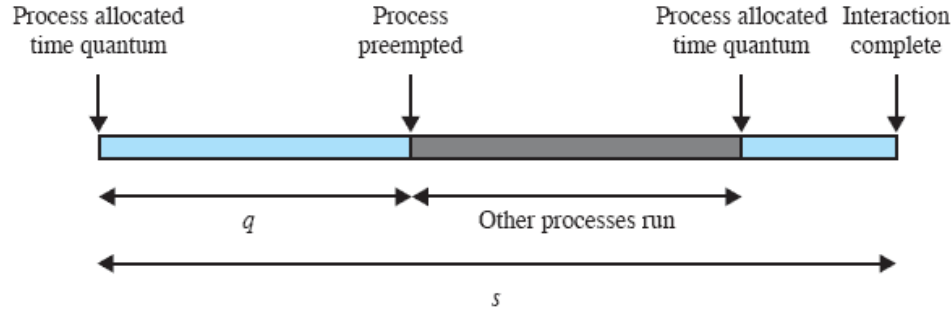


Effect of Size of Preemption Time Quantum



(a) Time quantum greater than typical interaction


Effect of Size of Preemption Time Quantum



(b) Time quantum less than typical interaction

Figure 9.6 Effect of Size of Preemption Time Quantum

'Virtual Round Robin'




Scheduling overview

- (Non) preemptive
 - Priority
 - FCFS
 - Round Robin
 - Shortest process Next
 - Shortest remaining time
 - Highest response ratio next
 - Feedback Sheduling
- } Prediction needed

Shortest Process Next

- Nonpreemptive policy
- Process with shortest expected processing time is selected next
- Short process jumps ahead of longer processes

Shortest Process
Next (SPN)



Shortest Process Next

Overall performance is significantly improved *throughput & waiting time*,

but:

- Predictability of longer processes is reduced
- If estimated time for process not correct, the operating system may abort it
- Possibility of starvation for longer processes

Calculating Program 'Burst' *based on observation of instances*

- Where:
 - T_i = processor execution time for the i th instance of this process
 - S_i = predicted value for the i th instance
 - S_1 = predicted value for first instance; not calculated

$$S_{n+1} = \frac{1}{n} \sum_{i=1}^n T_i$$


Exponential Averaging

- A common technique for predicting a future value on the basis of a time series of past values is exponential averaging

$$S_{n+1} = \alpha T_n + (1 - \alpha)S_n$$

Trust latest
measures
more

Exponential Smoothing Coefficients




$$S_{n+1} = \alpha T_n + (1 - \alpha)S_n$$

- Advantage low α ?
- Advantage high α ?
- What's ($\alpha \rightarrow 0$) ?


Figure 9.8 Exponential Smoothing Coefficients

Use Of Exponential Averaging



(a) Increasing function

Use Of Exponential Averaging



(b) Decreasing function


Exponential Averaging

Conclusions

- exponential averaging tracks changes in process behaviour faster than does simple averaging
- larger value of α results in a more rapid reaction to the change in the observed value
- Lower value of α better level out peak changes in the observed value

Shortest Remaining Time

- Preemptive version of **Shortest Process Next** policy
- Must estimate processing time and choose the shortest (+administer *remaining time*)



-C Preempts B
-Better than SPN
-Risk of starvation
of longer processes


Highest Response Ratio Next

- Problem in SPN: can result in starvation
- Choose next process with the largest ratio

$$\text{Ratio} = \frac{\text{time spent waiting} + \text{expected service time}}{\text{expected service time}}$$

$$\text{Ratio} \quad R = \frac{w + s}{s}$$

**Highest Response
Ratio Next (HRRN)**




Scheduling overview

- (Non) preemptive
 - Priority
 - FCFS
 - Round Robin
 - Shortest process Next
 - Shortest remaining time
 - Highest response ratio next
 - Feedback Sheduling
- } Prediction needed

Feedback Scheduling


We do not know length process!

- Penalize jobs that have been running longer (*~prefer shorter jobs*)
- Preemption (time q)
- Don't know remaining time process needs to execute



Feedback Performance

- Variations exist, simple version **pre-empts** periodically, similar to round robin
 - But can lead to **starvation** (e.g. lots of small jobs)





› FOR SOCIETY

Performance comparison

Normalized response time

- Any scheduling discipline that chooses the next item to be served **independent** of service time obeys the relationship:

$$\frac{T_r}{T_s} = \frac{1}{1 - \rho}$$

where

T_r = turnaround time or residence time; total time in system, waiting plus execution

T_s = average service time; average time spent in Running state

ρ = processor utilization


Experiment

In Practice:


- Arrival distribution:
 - Poisson
- Length per process
 - Exponential

Experiment:

- Priority = Length (2 classes)
Equal number long and short processes




Overall Normalized Response Time




- Preference priority for large jobs $r > 0.8$
- Better with preemption
- Almost insignificant

Normalized Response Time for (HP) Shorter Processes



Normalized Response Time for (LP) Longer Processes



Normalized Turnaround Time and Waiting Time

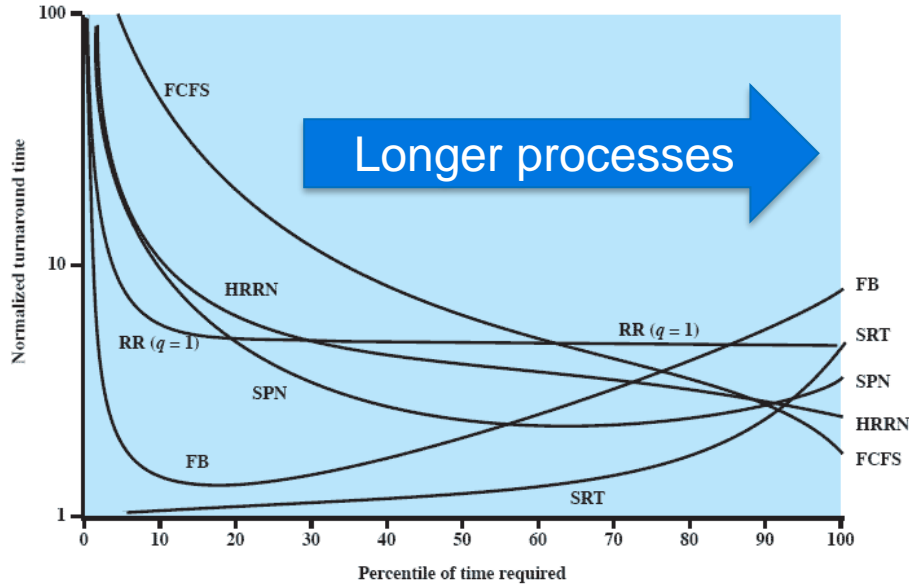


Figure 9.14 Simulation Results for Normalized Turnaround Time




Figure 9.15 Simulation Results for Waiting Time

Mini Assessment

Assignment

